

Inventory Problem of a Perishable Product
STA4001

GUO Yuanxin
118010084

November 3, 2019

Contents

1	Problem Formulation	3
1.1	Parameters in this problem	3
2	Analysis of the Process	4
2.1	Verification of Markovity	4
2.2	Transition Probabilities	4
2.3	Value function	5
3	Monte Carlo Method	5
4	Neural Network Approach	6
4.1	Subquestion (a)	7
4.2	Subquestion (b)(c)	7
A	Monte Carlo Python Code	9
B	Neural Network Python Code	10

1 Problem Formulation

Consider a perishable product with a maximum lifetime of L periods. Thus, at the end of period t , there may be leftover inventory with a remaining lifetime $1 \leq j \leq L - 1$. Let X_{tj} be the amount of inventory with remaining lifetime less or equal to j , then the system state in period t is $X_t = (X_{t1}, X_{t2}, \dots, X_{t,L-1})$.

At the beginning of period t , an order for $Q - X_{t-1,L-1}$ units of new inventory placed at the end of period $t - 1$ arrives, bringing the total inventory level to Q . A random demand D_t then occurs and is met immediately as much as possible with inventory from the oldest to the newest. Demand across different periods are assumed to be i.i.d. and follow Poisson distribution with mean λ . Unmet demand is lost. At the end of period t expired inventory is disposed at a cost c_d per unit, and leftover inventory incurs a holding cost h per unit.

We will soon show that $X = \{X_t, t = 0, 1, \dots\}$ is a DTMC. We use \mathcal{S} to denote the state space of the Markov chain. The objective is to compute the value function:

$$v(x) = \mathbb{E}[\sum_{n=0}^{\infty} \beta^n g(X_n) | X_0 = x]$$

for different initial states $x \in \mathcal{S}$, where $g(x)$ is the single period expected profit function given the state x in the previous period.

1.1 Parameters in this problem

Notation	Parameter	Value
c_p	Selling price	\$1.0
c_v	Variable cost	\$0.4
c_d	Disposal cost	\$0.1
h	Holding cost	\$0.1
λ	Demand rate	17
Q	Total inventory level at the beginning of period	90
L	Maximum lifetime of perishable product	7
β	Discount factor	0.8

2 Analysis of the Process

2.1 Verification of Markovity

To show that this process is a DTMC, we aim to find the relationship between X_n and X_{n+1} . If X_{n+1} is a function of X_n and a random variable, the Markovity is thus shown.

During each period, the system state undergoes following changes:

- a. The total inventory level is brought up to 90, while the remaining lifetime of each product reduce by 1. We use vector $Y_t = (Y_{t0}, Y_{t1}, \dots, Y_{t6})$ to denote the inventory state at the moment, where Y_{tj} be the amount of inventory with remaining lifetime less or equal to j . At the moment, $Y_t = (X_{t1}, X_{t2}, \dots, X_{t,L-1}, 90)$.
- b. The products are sold at a random demand D_t according to Poisson distribution from the oldest to the newest. Unmet demand is lost. $Y_t = ((X_{t1} - D_t)^+, (X_{t2} - D_t)^+, \dots, (X_{t6} - D_t)^+, (90 - D_t)^+)$.
- c. The expired inventory, namely the products with remaining lifetime equal to 0, are disposed. $Y_t = (0, (X_{t2} - D_t)^+ - (X_{t1} - D_t)^+, \dots, (X_{t6} - D_t)^+ - (X_{t1} - D_t)^+, (90 - D_t)^+ - (X_{t1} - D_t)^+)$.
- d. At the beginning of period $t+1$, the system state becomes $X_{t+1} = ((X_{t2} - D_t)^+ - (X_{t1} - D_t)^+, \dots, (X_{t6} - D_t)^+ - (X_{t1} - D_t)^+, (90 - D_t)^+ - (X_{t1} - D_t)^+)$.

Hence, it is not difficult to see that the expression of X_{t+1} contains solely terms of X_t and random variable D_t , thus proving that $X = \{X_t, t = 0, 1, \dots\}$ is a DTMC.

2.2 Transition Probabilities

Consider random variable $D_t \sim (\text{Poisson}, 17)$ with pmf

$$p_{D_t}(x) = \frac{17^x e^{-17}}{x!}.$$

Then the transition probability p_{ij} is:

$$\begin{cases} p_{D_t}(x), & j = ((i_2 - x)^+ - (i_1 - x)^+, \dots, (i_6 - x)^+ - (i_1 - x)^+, (90 - x)^+ - (i_1 - x)^+), \ x < 90 \\ \sum_{x=90}^{\infty} p_{D_t}(x), & j = ((i_2 - x)^+ - (i_1 - x)^+, \dots, (i_6 - x)^+ - (i_1 - x)^+, (90 - x)^+ - (i_1 - x)^+), \ x = 90 \\ 0, & \text{otherwise} \end{cases}$$

2.3 Value function

The value function $g(x)$ takes previous state X_t as the input and outputs the profit dependent on the realization of the random variable D_{t+1} . The analysis is analogous to that in (2.1).

- a. Bringing the total inventory up to 90 requires ordering $(90 - X_{t6})$ products.
This costs: $c_v \cdot (90 - X_{t6})$.
- b. The profit due to selling products depends on the outcome of the demand random variable. The number of sales is $\min(D_{t+1}, 90)$.
Hence this part of the profit is: $c_p \cdot \min(D_{t+1}, 90)$
- c. The items with lifetime equal to 0 are expired and to be disposed. There are $(X_{t1} - D_{t+1})^+$ of them.
The disposal cost is: $c_d \cdot (X_{t1} - D_{t+1})^+$.
- d. Finally, all of the rest inventory at a total quantity of $(90 - D_t)^+ - (X_{t1} - D_t)^+$ requires a holding cost.
The holding cost is: $h \cdot (90 - D_t)^+ - (X_{t1} - D_t)^+$.

In summary, the explicit form of the value function $g(x)$ is:

$$\begin{aligned} g(x) &= -c_v(90 - X_{t6}) + c_p \min(D_{t+1}, 90) - c_d(X_{t1} - D_{t+1})^+ - h [(90 - D_t)^+ - (X_{t1} - D_t)^+] \\ &= \min(D_{t+1}, 90) - 0.4(90 - X_{t6}) - 0.1(X_{t1} - D_{t+1})^+ - 0.1 [(90 - D_t)^+ - (X_{t1} - D_t)^+] \end{aligned}$$

3 Monte Carlo Method

To evaluate the expectation of the value function $v(x)$, it is almost impossible to compute analytically as the state space grows exponentially with L . Thus, we consider using simulation to approximate the result. A common approach is using the Monte Carlo method.

The basic idea of the Monte Carlo simulation is to find a set of sample paths and calculate $\mathbb{E}[\sum_{n=0}^{\infty} \beta^n g(X_n) | X_0 = x]$ over the set. The method approximates well provided both the time horizon and episode are large enough, justified by the Law of Large Numbers.

We run the code in (A) and obtain Figure 1 below. As the number of episodes and the length of time horizon increases, the estimation $\hat{v}(x)$ of the value function becomes more centered around 14.5. The fluctuations of both the estimation and the CI's become milder.

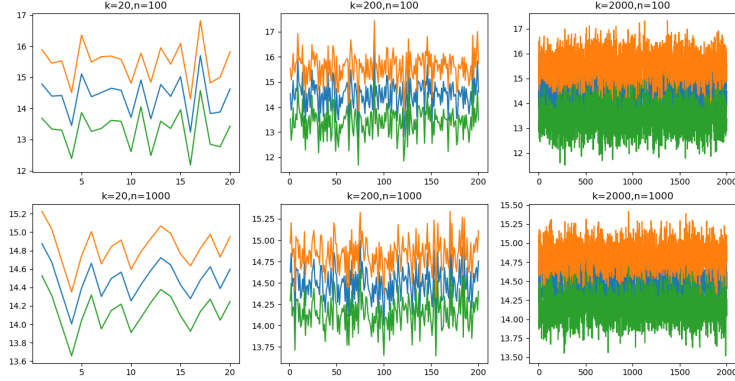


Figure 1: Monte Carlo Method.

4 Neural Network Approach

Consider $|\mathcal{S}|$, the size of the state space. We use the following Python code to compute it:

```
count = 0
for a in range(91):
    for b in range(a,91):
        for c in range(b,91):
            for d in range(c,91):
                for e in range(d,91):
                    for f in range(e,91):
                        count+=1
print("The size of the state space is",count)
```

The output is:

The size of the state space is 927048304

The state space is horribly large even for $L=7$, which makes it infeasible for the computer to store the $\hat{v}(x)$ for all $x \in \mathcal{S}$. Therefore, Monte Carlo method is only an expedient for a rather small subset of the state space.

One possible solution to estimate $v(x)$ for all $x \in \mathcal{S}$ is to train a neural network. We can apply the method in (3) to obtain a dataset (\bar{X}, \bar{Y}) for a relatively small subset \bar{X} , where $\bar{Y} = \{\hat{v}(\bar{x}_k) | \bar{x}_k \in \bar{X}, k = 1, \dots, K\}$, K is the size of the sample

space. Now we can use (\bar{X}, \bar{Y}) to train a neural network to approximate the value function over the whole state space.

We run the code in (B) in order to obtain the neural network model and solve the subquestions in Part 2.

4.1 Subquestion (a)

For $x = (10, 20, 30, 40, 50, 80)$, the estimated values are as follows:

$\tilde{v}(x)$ ($K = 20$)	14.52520264
$\tilde{v}(x)$ ($K = 200$)	14.41595547
$\tilde{v}(x)$ ($K = 2000$)	14.52172661
$\hat{v}(x)$	17.06012879080918
CI(95%)	[16.719012309096062, 17.4012452725223]

We discover that all of the predictions do not lie in the 95% confidence interval. A tentative explanation of this result is that the x value that we test is an outlier. In other words, the probability of visiting the state is scarce. The neural network model may not perform well when receiving an unlikely input.

4.2 Subquestion (b)(c)

Figure (2) includes the Monte Carlo estimation, neural network prediction and the confidence interval.




Figure 2: Neural Network Prediction.

The blue line stands for \hat{v} , the Monte Carlo estimation. The orange line and the green line are respectively the left and right endpoints of the confidence interval. The red, purple, and brown line correspond to the neural network prediction when $K = 20, 200, 2000$.

The mean square error for the test set is outputted below:

```
Mean squared error for K=20:  0.05379644691028611
Mean squared error for K=200:  1.8122860494687556
Mean squared error for K=2000: 0.5082348063789286
```


A Monte Carlo Python Code

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 x = [0,0,0,0,0,0]
5
6 x_bar_list = []
7 K = [20,200,2000]
8 beta = 0.8
9
10 k = [np.arange(1,21),np.arange(1,201),np.arange(1,2001)]
11
12 def transition(z,d):
13     y = z + [90]
14     for i in range(7):
15         y[i] = max(y[i]-d,0)
16     if y[0] > 0:
17         for j in range(1,7):
18             y[j] = y[j]-y[0]
19         y[0] = 0
20     return y[1:7]
21
22 def g(z,d):
23     disposal = max(z[0]-d,0)
24     refill = 90-z[5]
25     sale = min(d,90)
26     profit = 1*sale-0.4*refill-0.1*disposal-0.1*(90-sale-disposal)
27     return profit
28
29 fig, axs = plt.subplots(2,3)
30
31 for a in range(3):
32     counter = 0
33     while True:
34         r = np.random.poisson(17)
35         x = transition(x,r)
36         counter += 1
37         if counter % 1000 == 0:
38             x_bar_list.append(x)
39         if len(x_bar_list) == K[a]:
40             break
41
42     estimate_list_100 = []
43     estimate_list_1000 = []
44     std_list_100 = []
45     std_list_1000 = []
46     for state in x_bar_list:
```

```

47     x = state
48     profit_list = []
49     for episode in range(1000):
50         profit = 0
51         for T in range(50):
52             r = np.random.poisson(17)
53             profit = profit + (beta**T)*g(x,r)
54             x = transition(x,r)
55             profit_list.append(profit)
56         state_profit_100 = np.mean(profit_list[:100])
57         state_profit_1000 = np.mean(profit_list)
58         state_std_100 = np.std(profit_list[:100])
59         state_std_1000 = np.std(profit_list)
60         estimate_list_100.append(state_profit_100)
61         estimate_list_1000.append(state_profit_1000)
62         std_list_100.append(state_std_100)
63         std_list_1000.append(state_std_1000)
64     std_list_100 = np.array(std_list_100)
65     std_list_1000 = np.array(std_list_1000)
66
67     axs[0,a].plot(k[a], estimate_list_100)
68     axs[0,a].plot(k[a], estimate_list_100+1.96*std_list_100/10)
69     axs[0,a].plot(k[a], estimate_list_100-1.96*std_list_100/10)
70     axs[0,a].set_title('k='+str(K[a])+',n=100')
71     axs[1,a].plot(k[a], estimate_list_1000)
72     axs[1,a].plot(k[a], estimate_list_1000+1.96*std_list_1000/np.
73                    sqrt(1000))
74     axs[1,a].plot(k[a], estimate_list_1000-1.96*std_list_1000/np.
75                    sqrt(1000))
76     axs[1,a].set_title('k='+str(K[a])+',n=1000')
77 plt.show()

```

B Neural Network Python Code

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3  from sklearn.metrics import mean_squared_error
4  from keras.models import Sequential
5  from keras.layers import Dense
6
7  class MyNeuralNetwork:
8      def __init__(self, input_dimension):
9          # define the keras model
10         self.model = Sequential()
11         self.model.add(Dense(10, input_dim=input_dimension,

```

```

12         activation='relu'))
13     self.model.add(Dense(6, activation='relu'))
14     self.model.add(Dense(1, activation='linear'))
15     # compile the keras model
16     self.model.compile(loss='mean_squared_error', optimizer='
        adam')
17
18 class Scaler:
19     # save mean and variance of x, y sets
20     def __init__(self, x, y):
21         self.x_mean = np.mean(x, axis=0)
22         self.y_mean = np.mean(y, axis=0)
23         self.x_std = np.std(x, axis=0)
24         self.y_std = np.std(y, axis=0)
25
26
27     def get_x(self):
28         # return saved mean and variance of x
29         return self.x_std, self.x_mean
30
31     def get_y(self):
32         # return saved mean and variance of y
33         return self.y_std, self.y_mean
34
35     def transition(z,d):
36         y = z + [90]
37         for i in range(7):
38             y[i] = max(y[i]-d,0)
39         if y[0] > 0:
40             for j in range(1,7):
41                 y[j] = y[j]-y[0]
42             y[0] = 0
43         return y[1:7]
44
45     def g(z,d):
46         disposal = max(z[0]-d,0)
47         refill = 90-z[5]
48         sale = min(d,90)
49         profit = 1*sale-0.4*refill-0.1*disposal-0.1*(90-sale-disposal)
50         return profit
51
52     def data_generator(x,K):
53         counter = 0
54         x_bar_list = []
55         beta = 0.8
56         while True:
57             r = np.random.poisson(17)
58             x = transition(x,r)

```

```

59         counter += 1
60         if counter == 1000:
61             x_data = x
62             x_bar_list.append(x)
63         elif (counter % 1000 == 0) and (x not in x_bar_list):
64             x_data = np.vstack([x_data,x])
65             x_bar_list.append(x)
66         if len(x_bar_list) == K:
67             break
68
69     estimate_list_100 = []
70     std_list_100 = []
71     for state in x_bar_list:
72         x = state
73         profit_list = []
74         for episode in range(1000):
75             profit = 0
76             for T in range(50):
77                 r = np.random.poisson(17)
78                 profit = profit + (beta**T)*g(x,r)
79                 x = transition(x,r)
80             profit_list.append(profit)
81         state_profit_100 = np.mean(profit_list)
82         state_profit_std_100 = np.std(profit_list)
83         estimate_list_100.append(state_profit_100)
84         std_list_100.append(state_profit_std_100)
85     estimate_list_100 = np.array(estimate_list_100)
86     std_list_100 = np.array(std_list_100)
87     return x_data, estimate_list_100[:, np.newaxis], std_list_100
88        [:, np.newaxis]
89
90
91     beta = 0.8
92     initial = [0,0,0,0,0,0]
93     x_1, y_1 = data_generator(initial,20)[:2]
94     x_2, y_2 = data_generator(initial,200)[:2]
95     x_3, y_3 = data_generator(initial,2000)[:2]
96     neural_network_1 = MyNeuralNetwork(input_dimension=6)
97     neural_network_2 = MyNeuralNetwork(input_dimension=6)
98     neural_network_3 = MyNeuralNetwork(input_dimension=6)
99
100
101     normalizer_1 = Scaler(x_1, y_1)
102     std_x_1, mean_x_1 = normalizer_1.get_x()
103     x_1_norm = (x_1 - mean_x_1) / (std_x_1+0.000001)
104     std_y_1, mean_y_1 = normalizer_1.get_y()
105     y_1_norm = (y_1 - mean_y_1) / (std_y_1+0.000001)

```

```

106
107 neural_network_1.model.fit(x_1_norm, y_1_norm, epochs=100,
    batch_size=8)
108
109
110 normalizer_2 = Scaler(x_2, y_2)
111 std_x_2, mean_x_2 = normalizer_2.get_x()
112 x_2_norm = (x_2 - mean_x_2) / (std_x_2+0.000001)
113 std_y_2, mean_y_2 = normalizer_2.get_y()
114 y_2_norm = (y_2 - mean_y_2) / (std_y_2+0.000001)
115
116 neural_network_2.model.fit(x_2_norm, y_2_norm, epochs=100,
    batch_size=8)
117
118
119 normalizer_3 = Scaler(x_3, y_3)
120 std_x_3, mean_x_3 = normalizer_3.get_x()
121 x_3_norm = (x_3 - mean_x_3) / (std_x_3+0.000001)
122 std_y_3, mean_y_3 = normalizer_3.get_y()
123 y_3_norm = (y_3 - mean_y_3) / (std_y_3+0.000001)
124
125 neural_network_3.model.fit(x_3_norm, y_3_norm, epochs=100,
    batch_size=8)
126
127
128 x_0 = np.array([10,20,30,40,50,80])
129 x_prime = np.vstack([x_0, x_0])
130 x_0_norm_1 = (x_prime - mean_x_1) / (std_x_1+0.000001)
131 x_0_norm_2 = (x_prime - mean_x_2) / (std_x_2+0.000001)
132 x_0_norm_3 = (x_prime - mean_x_3) / (std_x_3+0.000001)
133 v_tilde_norm_1 = neural_network_1.model.predict(x_0_norm_1)
134 v_prime = v_tilde_norm_1 * std_y_1 + mean_y_1
135 v_tilde_1 = v_prime[0]
136 v_tilde_norm_2 = neural_network_2.model.predict(x_0_norm_2)
137 v_prime = v_tilde_norm_2 * std_y_2 + mean_y_2
138 v_tilde_2 = v_prime[0]
139 v_tilde_norm_3 = neural_network_3.model.predict(x_0_norm_3)
140 v_prime = v_tilde_norm_3 * std_y_3 + mean_y_3
141 v_tilde_3 = v_prime[0]
142
143
144 profit_list = []
145 for episode in range(1000):
146     profit = 0
147     x = x_0.tolist()
148     for T in range(50):
149         r = np.random.poisson(17)
150         profit = profit + (beta**T)*g(x,r)

```

```

151         x = transition(x,r)
152         profit_list.append(profit)
153
154     v_hat = np.mean(profit_list)
155     v_sigma = np.std(profit_list)
156     v_CI_lower = v_hat-1.96*v_sigma/np.sqrt(1000)
157     v_CI_upper = v_hat+1.96*v_sigma/np.sqrt(1000)
158
159
160     print("v_tilde_1:",v_tilde_1)
161     print("v_tilde_2:",v_tilde_2)
162     print("v_tilde_3:",v_tilde_3)
163     print("v_hat:",v_hat)
164     print("CI (95%):["+str(v_CI_lower)+", "+str(v_CI_upper)+"]")
165
166     x_test, y_test, std_test = data_generator([0,0,0,0,0,0],50)
167     x_test_norm_1 = (x_test - mean_x_1) / (std_x_1+0.000001)
168     x_test_norm_2 = (x_test - mean_x_2) / (std_x_2+0.000001)
169     x_test_norm_3 = (x_test - mean_x_3) / (std_x_3+0.000001)
170     y_pred_1 = neural_network_1.model.predict(x_test) * std_y_1 +
        mean_y_1
171     y_pred_2 = neural_network_2.model.predict(x_test) * std_y_2 +
        mean_y_2
172     y_pred_3 = neural_network_3.model.predict(x_test) * std_y_3 +
        mean_y_3
173
174     mse_1 = mean_squared_error(y_test, y_pred_1)
175     mse_2 = mean_squared_error(y_test, y_pred_2)
176     mse_3 = mean_squared_error(y_test, y_pred_3)
177     print('Mean_squared_error_for_K=20:', mse_1)
178     print('Mean_squared_error_for_K=200:', mse_2)
179     print('Mean_squared_error_for_K=2000:', mse_3)
180
181
182     plt.plot(y_test, label="v_hat", lw=2)
183     plt.plot(y_test-1.96*std_test/np.sqrt(1000), label="CI_lower", lw
        =2)
184     plt.plot(y_test+1.96*std_test/np.sqrt(1000), label="CI_upper", lw
        =2)
185     plt.plot(y_pred_1, label="v_tilde_1", lw=2)
186     plt.plot(y_pred_2, label="v_tilde_2", lw=2)
187     plt.plot(y_pred_3, label="v_tilde_3", lw=2)
188     plt.legend()
189     plt.show()

```